

Predicting Growth and Remodeling of Cardiovascular Tissue

change in mass

change in microstructure

Dr.-Ing. Martin R. Pfaller

Assistant Professor

Department of Biomedical Engineering, Yale University

biomechanics.yale.edu

With much inspiration from Drs. Amadeus Gebauer, Jonas Eichinger, and Jay Humphrey

Teacher Version

Yale

A Brief Introduction to Mechanobiology of Blood Vessels

The Origin of Mechanobiology

Davis' law, Henry Gassett Davis (1807–1896, USA), 1867:


“Ligaments, or any soft tissue, when put under even a moderate degree of tension, if that tension is unremitting, will elongate by the addition of new material; [...].”

Wolff's law, Julius Wolff (1836–1902, Germany), 1892 (*Das Gesetz der Transformation der Knochen*, the law of bone transformation):

“im Gefolge primärer Abänderungen der Form und Inanspruchnahme oder auch bloß der Inanspruchnahme der Knochen [sich] bestimmte, nach mathematischen Regeln eintretende Umwandlungen der inneren Architektur und ebenso bestimmte, denselben mathematischen Regeln folgende sekundäre Umwandlungen der äußeren Form der betreffenden Knochen vollziehen.”

- bone mass proportional to load (higher / lower load $\Rightarrow$ growth / atrophy)
- first *mathematical* relation between load and growth & remodeling in biological tissue


Vascular tissue has a preferred mechanical state



Tension T varies by a **factor of 26**

$$T = P a$$



T and σ_θ





Tension / lamella varies by a **factor of 2.8**

$$\frac{T}{h} = \frac{P a}{h} = \sigma_\theta$$


Mechanical Homeostasis



(Mal)adaptation




Mechanobiological Regulation



An Idealized Blood Vessel

A vessel is a **pressurised, perfused tube**. Two loads act on the wall: **blood pressure** P pushes it outward, and **blood flow** Q drags its inner surface. Cells sense both.



a inner radius, h wall thickness

P lumen pressure $\rightarrow$ **circumferential stress** $\sigma_\theta \sim 100$ kPa

Q volumetric flow $\rightarrow$ **wall shear stress** $\tau_w \sim 1$ Pa

Two Loads, Two Formulae, Two Adaptations

Pressure P sets the **hoop stress**, flow Q sets the **wall shear stress**.

Pressure (Laplace)

$$\sigma_{\theta} = \frac{Pa}{h} \quad (1)$$

Flow (Poiseuille)

$$\tau_w = \frac{4\mu Q}{\pi a^3} \quad (2)$$

✂ Your turn: Pressure adaptation


Question. Pressure **doubles**. How does the wall thickness change? $\sigma_\theta = \frac{Pa}{h}$

✓ Answer: Pressure adaptation

$P'/h' = P/h$ with $P' = 2P$ gives

$$h' = \left(\frac{P'}{P}\right) h = 2h. \quad (3)$$

Wall thickens *in proportion* to pressure.



$P \uparrow$

✂ Your turn: Flow adaptation


Question. Flow rises **8-fold**. How does the radius change? $\tau_w = \frac{4\mu Q}{\pi a^3}$

✓ Answer: Flow adaptation

$(a'/a)^3 = Q'/Q = 8$, so

$$a' = \left(\frac{Q'}{Q}\right)^{1/3} a = 2a. \quad (4)$$

Reading. Radius grows as the *cube root* of flow.




$Q \uparrow$

Finite-Strain Fundamentals


Reference and Current Configuration

A body is described by a **reference** $\mathcal{B}_0$ (points $\mathbf{X}$) and a **current** $\mathcal{B}_t$ (points $\mathbf{x}$) configuration, related by the motion $\mathbf{x} = \varphi(t, \mathbf{X})$.

3D 3D geometric model



0D 0D reduced model



The Deformation Gradient

The **deformation gradient** maps reference line elements to current ones.

3D a 3×3 tensor

$$d\mathbf{x} = \mathbf{F} d\mathbf{X}, \quad \mathbf{F} = \frac{\partial \boldsymbol{\varphi}}{\partial \mathbf{X}}, \quad J = \det \mathbf{F} \quad (5)$$

J is the volume ratio, **incompressibility** means $J = 1$.

0D a scalar

$$\mathbf{F} \longrightarrow \lambda = \frac{\text{current length}}{\text{reference length}}, \quad \lambda = 1 \text{ undeformed.}$$

Strain: A Rotation-Free Measure of Stretch

Removing rigid rotation gives the **right Cauchy–Green** tensor and the **Green–Lagrange strain**.

3D symmetric 3×3 tensors

$$\mathbf{C} = \mathbf{F}^T \mathbf{F} = (\mathbf{R}\mathbf{U})^T \mathbf{R}\mathbf{U} = \mathbf{U}^T \mathbf{R}^T \mathbf{R} \mathbf{U} = \mathbf{U}^2, \quad \mathbf{E} = \frac{1}{2}(\mathbf{C} - \mathbf{I}) \quad (6)$$

0D scalars

$$C = \lambda^2, \quad E = \frac{1}{2}(\lambda^2 - 1).$$

Stress: "Force per Area"

Cauchy stress: *current* force per unit *current* area.

3D spatial symmetric 3×3 tensor (from 2nd Piola–Kirchhoff stress S)

$$\boldsymbol{\sigma} = \frac{1}{J} \mathbf{F} \mathbf{S} \mathbf{F}^T, \quad \mathbf{S} = 2 \frac{\partial W}{\partial \mathbf{C}}, \quad W \text{ strain energy density function.}$$

0D scalar reduction ($C = \lambda^2$, $J = 1$), the σ_θ in Laplace (1)

$$\sigma = \frac{1}{J} \lambda S \lambda = \lambda^2 S = 2\lambda^2 \frac{\partial W}{\partial C} = \lambda \frac{dW}{d\lambda}.$$

Hyperelasticity: Stress from a Stored-Energy Function

A **hyperelastic** material stores energy W (deformation is path-independent).

3D energy $W(\mathbf{C})$, stress $\mathbf{S} = 2 \partial W / \partial \mathbf{C}$

$$W_{\text{neo-Hooke}} = \frac{c}{2} (I_1 - 3), \quad W_{\text{Fung}} = \frac{c_1}{4c_2} \left[e^{c_2(I_1-1)^2} - 1 \right]$$

$I_1 = \text{tr } \mathbf{C}$ (isotropic), $I_4 = \mathbf{f} \cdot \mathbf{C} \mathbf{f} = \lambda_f^2$, the squared stretch along the reference fiber direction $\mathbf{f}$ (anisotropic).

0D reduction used here, one stretch λ , scalar $\sigma = \lambda \, dW/d\lambda$

$$\sigma_{\text{neo-Hooke}} = c(\lambda^2 - \lambda^{-1}), \quad \sigma_{\text{Fung}} = c_1 \lambda^2 (\lambda^2 - 1) e^{c_2(\lambda^2-1)^2} \quad (7)$$

Kinematic Growth

Multiplicative Decomposition: Growth vs. Elasticity

Split the deformation into a stress-free **growth** part F_g and a stress-carrying **elastic** part F_e .

3D a product of tensors


$$\mathbf{F} = \mathbf{F}_e \mathbf{F}_g, \quad \text{only } \mathbf{F}_e \text{ stores energy.}$$

0D scalar stretches

$$\lambda = \lambda_e \lambda_g \tag{8}$$

Multiplicative Decomposition: Growth vs. Elasticity

F_g yields an *incompatible* stress-free state, F_e restores compability via residual stress.



Kinematic Growth Tensor

$\mathbf{F}_g$ encodes *how* the tissue is assumed to add mass. A common choice grows the tissue an amount θ along a fixed direction $\mathbf{f}$:

3D prescribed growth tensor

$$\mathbf{F}_g = \mathbf{I} + (\theta - 1) \mathbf{f} \otimes \mathbf{f} = \text{diag}(1, 1, \theta) \quad (9)$$

0D growth stretch

$$\lambda_g \text{ (scalar growth stretch),} \quad \theta = \det \mathbf{F}_g = M/M_0.$$

Growth Law

The growth *pattern* $\mathbf{F}_g$ is typically fixed. Its rate $\dot{\mathbf{F}}_g = \dot{\theta} \mathbf{f} \otimes \mathbf{f}$ reduces to *one* scalar ODE:

$$\frac{d\theta}{dt} = k_g \theta \left(\frac{\bar{\sigma}}{\bar{\sigma}_h} - 1 \right) \quad (10)$$

Since $k_g \theta > 0$, the sign of the response follows the stress deviation:

$$\frac{d\theta}{dt} \begin{cases} > 0 & \bar{\sigma} > \bar{\sigma}_h & \text{(growth)} \\ = 0 & \bar{\sigma} = \bar{\sigma}_h & \text{(homeostasis)} \\ < 0 & \bar{\sigma} < \bar{\sigma}_h & \text{(atrophy)} \end{cases}$$

Getting Started with the Code

One runner drives *every* exercise: you never edit Python, just point it at a small input file. **First-time setup** (needs Python ≥ 3.10 ; dependencies: NumPy, SciPy, Matplotlib, PyYAML):

```
git clone https://github.com/YaleCBL-Teaching/Growth-Remodeling
cd Growth-Remodeling
```

uv (recommended):
uv sync

pip (your own Python):
python -m venv .venv
pip install -e .

What is where:

- `run.py`: the single runner for all exercises
- `configs/`: one YAML input file per exercise (*you edit these*)
- `src/gr/`: the models & maths (read, don't edit)
- `docs/`: documentation
- `slides/`: semianr notes

⚙️ Setup: Example (hypertensive artery, ex01)

Example. A pressurised artery is made *hypertensive* by a sustained pressure step. Kinematic growth thickens the wall. Does the wall stress $\bar{\sigma}/\bar{\sigma}_h$ return to the *prescribed* set-point?

Running an Exercise

Point the runner at the input file (`uv run`, or plain `python` once the `venv` is activated):

```
uv run python run.py configs/ex01_kinematic_growth.yaml
```

You get: the key numbers ($\bar{\sigma}/\bar{\sigma}_h$, mass, diverged) print to the terminal, and the time history opens as a plot.

Change one value: edit the YAML, or override on the command line:

```
uv run python run.py configs/ex01_kinematic_growth.yaml -set insult.pressure_factor=2
```

Parametric study: one curve per value:

```
... -sweep insult.pressure_factor=1.5,2,3
```




Your turn: Kinematic growth

Question. Does the wall stress return to the prescribed set-point? Try larger pressure steps. Does the outcome ever change?

✓ Answer: Kinematic growth


The stress plateau *always* lands on the prescribed set-point. The wall simply thickens more for a larger pressure step. Kinematic growth *imposes* stability.

Kinematic growth of the artery under hypertension (P: 1.0 $\rightarrow$ 1.5x)




The Constrained Mixture Model

One Framework, Many Applications



Ubiquitous Vascular G&R

The Idea: Track Constituents Over Time



κ : configuration • τ : deposition time • s : current time • α : constituent • G^α : deposition stretch • $F_{n(\tau)}^\alpha$: constituent deformation

Load-Bearing Vessel Constituents

The wall is a mixture of three constituents, each with a reference **mass fraction** ϕ_0^α ($\sum_\alpha \phi_0^\alpha = 1$) and a **deposition stretch** G^α , the prestretch cells lay it down at, which *defines* its homeostatic stress $\sigma_h^\alpha = \sigma^\alpha(G^\alpha)$; the mixture set-point is $\bar{\sigma}_h = \sum_\alpha \phi_0^\alpha \sigma_h^\alpha$.

Constituent	Role	ϕ_0	G^α	law; turnover $1/k_d$
Elastin	bears load at low stretch	0.30	1.40	neo-Hooke; none (decades)
Collagen	stiff, anti-rupture	0.35	1.08	Fung; fast (~ 80 d)
Smooth muscle	passive + active	0.35	1.20	Fung; fast (~ 80 d)

Elastin cannot be renewed ($k_d^e = 0$), its loss drives *aneurysm*. Collagen and SMC **turn over**, so the wall can *adapt*.

Every Cohort Shares One Stretch

A cohort of α deposited at τ is laid down at G^α and then stretches *with the mixture*: every constituent and cohort is **constrained** to the same tissue stretch $\lambda(s)$.

3D constituent-specific deformation

$$\mathbf{F}_{n(\tau)}^\alpha(s) = \mathbf{F}(s) \mathbf{F}^{-1}(\tau) \mathbf{G}^\alpha(\tau).$$

0D scalar elastic stretch

$$\lambda_{n(\tau)}^\alpha(s, \tau) = \frac{\lambda(s)}{\lambda(\tau)} G^\alpha \quad (11)$$

Mass Balance

Track each constituent's **apparent mass density** ρ^α per current mixture volume, $\rho = \sum_\alpha \rho^\alpha$.

Spatial mass balance (one per constituent):

$$\frac{\partial \rho^\alpha}{\partial s} + \underbrace{\operatorname{div}(\rho^\alpha \mathbf{v}^\alpha)}_{\text{transport}} = \underbrace{\bar{m}^\alpha}_{\text{net mass exchange}}, \quad \alpha = 1, \dots, N.$$

The net rate is *signed*:

$$\bar{m}^\alpha = \underbrace{m^\alpha}_{\text{production} \geq 0} - \underbrace{r^\alpha}_{\text{removal} \geq 0} \quad \left\{ \begin{array}{ll} < 0 & \text{atrophy} \\ = 0 & \text{maintenance} \\ > 0 & \text{growth} \end{array} \right.$$

From Balance to a Heredity Integral

Assumption 1 (constrained mixture): every constituent moves with the mixture, $v^\alpha = v$.

Assumption 2 (prescribed rates): a true production rate $m^\alpha(\tau) \geq 0$ and a survival fraction $q^\alpha(s, \tau) \in [0, 1]$. Integrating from the distant past to s :

$$\rho^\alpha(s) = \underbrace{\rho^\alpha(0) Q^\alpha(s)}_{\text{survivors of original material}} + \underbrace{\int_0^s m^\alpha(\tau) q^\alpha(s, \tau) d\tau}_{\text{deposited at } \tau, \text{ still present at } s} \quad (12)$$

with $Q^\alpha(0) = 1$. **Assumption 3:** bulk density is nearly constant, $\rho(s) = \rho(0)$, mass changes track volume changes.

Linear Momentum Balance

A rule-of-mixtures for the stresses (no body forces) satisfies the classical linear-momentum balance:

$$\operatorname{div} \boldsymbol{\sigma} = \rho \boldsymbol{a} \xrightarrow{\text{quasi-static}} \operatorname{div} \boldsymbol{\sigma} = \mathbf{0}.$$

Quasi-static assumption: G&R evolves slowly at every time s , so the inertia $\rho \boldsymbol{a}$ is negligible.

Constitutive Relations: Cauchy Stress

In finite elasticity the Cauchy stress follows from a stored-energy function W (per unit reference mixture volume):

$$\boldsymbol{\sigma}(s) = \frac{2}{J} \mathbf{F}(s) \frac{\partial W(s)}{\partial \mathbf{C}(s)} \mathbf{F}^T(s), \quad \mathbf{C} = \mathbf{F}^T \mathbf{F}, \quad J = \det \mathbf{F}.$$

Rule of mixtures: $W = \sum_{\alpha} \hat{W}^{\alpha}$. Closing the model needs **three constituent-specific relations**:

- a *stored energy* W^{α} ,
- a *production rate* m^{α} , and
- a *removal* (survival) q^{α} .

Three Constituent-Specific Constitutive Relations

- 1. Material:** constituent stored energy W^α (e.g., neo-Hookean, Fung-type)
- 2. Rate of production:** up with wall stress and inflammation, down with wall shear:

$$m^\alpha(\tau) = m_o^\alpha \left(1 + \underbrace{\frac{K_\sigma^\alpha \Delta\sigma}{\text{circumferential stress}}}_{\text{circumferential stress}} - \underbrace{\frac{K_{\tau_w}^\alpha \Delta\tau_w}{\text{wall shear stress}}}_{\text{wall shear stress}} + \underbrace{\frac{K_\varphi^\alpha \Delta\rho_\varphi}{\text{inflammation}}}_{\text{inflammation}} \right) \geq 0 \quad (13)$$

- 3. Rate of removal:** a survival fraction that decays faster under stress perturbation:

$$q^\alpha(s, \tau) = \exp\left(-\int_\tau^s k_d^\alpha (1 + \omega \Delta\sigma^2(t) + \dots) dt\right) \in [0, 1] \quad (14)$$

Constituent Stored-Energy Functions

Each cohort carries its *own* stored energy W^α . Summing over cohorts with the *same* heredity structure gives the apparent (mixture-level) energy:

$$\rho \hat{W}^\alpha(s) = \underbrace{\rho^\alpha(0) Q^\alpha(s) W^\alpha(\mathbf{F}_{n(0)}^\alpha(s))}_{\text{original material}} + \underbrace{\int_0^s m^\alpha(\tau) q^\alpha(s, \tau) W^\alpha(\mathbf{F}_{n(\tau)}^\alpha(s)) d\tau}_{\text{later cohorts}}, \quad (15)$$

with elastic deformation

$$\mathbf{F}_{n(\tau)}^\alpha(s) = \mathbf{F}(s) \mathbf{F}^{-1}(\tau) \mathbf{G}^\alpha(\tau). \quad (16)$$

The mixture energy $W = \sum_\alpha \hat{W}^\alpha$ closes the mixture stress.

Special Cases: Before Perturbation and Homeostasis

(I) At $s = 0$ (just before a perturbation) only original material is present, and $Q^\alpha(0) = 1$:

$$\rho \hat{W}^\alpha(0) = \rho^\alpha(0) W^\alpha(\mathbf{F}_{n(0)}^\alpha(0)) \Rightarrow \hat{W}^\alpha(0) = \underbrace{\frac{\rho^\alpha(0)}{\rho(0)}}_{\phi^\alpha(0)} W^\alpha(0),$$

recovering the rule of mixtures $W(0) = \sum_\alpha \phi^\alpha(0) W^\alpha(0)$.

(II) **Homeostasis:** turnover in an *unchanging* configuration, so $\mathbf{F}_{n(\tau)}^\alpha(s) = \mathbf{F}_{n(0)}^\alpha(s)$ (independent of τ). Then W^α factors out and, using the mass balance,

$$\rho \hat{W}^\alpha(s) = \underbrace{\left(\rho^\alpha(0) Q^\alpha(s) + \int_0^s m^\alpha q^\alpha d\tau \right)}_{= \rho^\alpha(s)} W^\alpha(\mathbf{F}_{n(0)}^\alpha(s)) \Rightarrow \hat{W}^\alpha(s) = \phi^\alpha(s) W^\alpha(s). \quad (17)$$

Perfect Maintenance

Take *constant* rates: $m^\alpha = m_0$, $Q^\alpha(s) = e^{-K^\alpha s}$, $q^\alpha(s, \tau) = e^{-K^\alpha(s-\tau)}$. The mass balance integrates to

$$\rho^\alpha(s) = \underbrace{\rho^\alpha(0) e^{-K^\alpha s}}_{\text{original decays}} + \underbrace{\frac{m_0}{K^\alpha} (1 - e^{-K^\alpha s})}_{\text{production accumulates}}.$$

A *time-invariant* density ($\rho^\alpha(s) = \rho^\alpha(0)$ for all s) requires the two to balance exactly:

$$\text{perfect maintenance} \iff \frac{m_0}{K^\alpha} = \rho^\alpha(0). \quad (18)$$

Production m_0 then exactly offsets removal $K^\alpha \rho^\alpha$ (continuous turnover, unchanging tissue).

✂ Your turn: Computational Cost?

Question. At each of N time steps the mixture stress sums over *all* cohorts deposited so far. How does the total cost scale with N ?

Hints. Step k re-sums $\sim k$ stored cohorts; add up $\sum_{k=1}^N k$.

✓ Answer: Computational Cost?

Step k re-sums the $\sim k$ cohorts stored so far, so the total work is

$$\sum_{k=1}^N k = \frac{N(N+1)}{2} = \mathcal{O}(N^2)$$

quadratic in the number of time steps. The *homogenized* model (next) removes this history sum, cutting the cost to $\mathcal{O}(N)$.

⚙️ Setup: Example 2: Turnover under Hypertension (ex02)

Example. The same artery, now with collagen & SMC turning over (elastin does not). Which constituent grows to carry the load, and does turnover restore the set-point? Equilibrium each step: $\bar{\sigma}(\lambda) = \bar{\sigma}_h \gamma \lambda^2 / m$ (load factor γ , mass ratio m).

Run it.

```
uv run python run.py configs/ex02_constrained_mixture.yaml
```

Sets each constituent's `turnover_days` ($1/k_d$), gain (K_σ), and the `insult`, hypertension (γ) or aneurysm (`elastin_surviving`).

Parametric study.

```
... -set model.constituents.collagen.gain=3.0
```




Your turn: Constrained Mixture Model

Question. Does the same tissue *adapt* to hypertension but *run away* on the aneurysm insult? Which constituent grows to carry the load elastin used to bear?

✓ Answer: Constrained Mixture Model

Mild hypertension **adapts**: collagen and SMC grow, tissue stress returns to homeostatic. A severe elastin insult can **run away**. Faster turnover or higher gain reach the new state sooner.


Full constrained-mixture model: adaptation to hypertension (artery)



Homogenized CMM

The Mechanical Analog: A Maxwell Fluid with a Motor

Each constituent behaves like a **Maxwell element in parallel with a motor**: turnover relaxes stress toward a *prestress*, not zero.



σ^α	constituent stress (the load)
$\mathbb{C}^\alpha$	elastic stiffness (spring)
$\dot{\varrho}_0^\alpha / \varrho_0^\alpha + 1/T^\alpha$	turnover + degradation (dashpot)
$\bar{\sigma}_{\text{pre}}^\alpha$	deposition prestress (motor)

Back to the full CMM (0D):

- spring $\leftrightarrow$ constituent law $\sigma^\alpha(\lambda_e^\alpha)$
- dashpot $\leftrightarrow$ mass turnover / removal k_d^α
- motor $\leftrightarrow$ deposition stretch G^α

Growth (Mass) and Remodeling (Natural Configuration)

Two 0D ODEs per constituent replace the cohort history.

Growth: mass production $\propto$ mass, driven by the tissue-stress deviation.

$$\frac{dM^\alpha}{dt} = M^\alpha K_\sigma^\alpha k_d^\alpha \left(\frac{\bar{\sigma}}{\bar{\sigma}_h} - 1 \right) \quad (19)$$

Remodeling: new mass enters at natural stretch λ/G^α , old mass leaves at the current λ_n^α (elastic stretch $\lambda_e^\alpha = \lambda/\lambda_n^\alpha$).

$$\frac{d\lambda_n^\alpha}{dt} = \frac{m^\alpha}{M^\alpha} \left(\frac{\lambda}{G^\alpha} - \lambda_n^\alpha \right) \quad (20)$$

with production $m^\alpha = k_d^\alpha M^\alpha \Upsilon^\alpha$ and stimulus $\Upsilon^\alpha = 1 + K_\sigma^\alpha (\bar{\sigma}/\bar{\sigma}_h - 1)$. Holding λ fixed, remodeling relaxes each constituent's stress to its prestress σ_h^α (time constant $\sim 1/k_d^\alpha$), the Maxwell-with-motor behaviour.

✂ Your turn: Why Is the Homogenized Model Cheap?

Question. How does the cost of the homogenized constrained mixture model scale with the number of steps N , compared with the full model?

✓ Answer: Why Is the Homogenized Model Cheap?

The state is *fixed-size*, two numbers $(M^\alpha, \lambda_n^\alpha)$ per constituent, and each step just advances two ODEs, with *no* history to re-sum:

$$\text{homogenized: } \mathcal{O}(N) \quad \text{vs.} \quad \text{full CMM: } \mathcal{O}(N^2)$$

Same time trajectory, far cheaper, and the fixed state maps onto a viscoelastic (Maxwell-with-motor) material, so it drops straight into standard finite-element solvers.

✂ Your turn: When do homogenized and full agree?

Question. The homogenized model keeps only the *mean* natural configuration per constituent. Under what assumptions should it match the full deposition-history model and when would you expect them to diverge?

✓ Answer: When do homogenized and full agree?

They agree when each constituent's cohorts share *nearly the same* natural configuration i.e. slow, gradual loading where deposition stretch G^{α} and turnover are roughly stationary, so the cohort spread stays narrow and the mean represents it well. **They diverge** under fast or strongly time-varying insults, when cohorts deposited at very different states coexist and a single mean stretch can no longer capture the distribution.

⚙️ Setup: Example (homogenized vs. full, ex03)

Example. Run both mixture theories on the *same* artery and insult. How closely does the homogenized model track the full model, and how much faster is it?

Run it.

```
uv run python run.py configs/ex03_homogenized_vs_full.yaml
```

Here study: compare overlays constrained_mixture, homogenized_cmm and kinematic_growth on one scenario, and prints each runtime and the mass-ratio gap between the first two (full vs. homogenized).

Parametric study. Stress the comparison with a harsher, longer insult:

```
... -set insult.elastin_surviving=0.3 -set insult.ramp=10
... -set simulate.t_end=4000      # does the agreement hold?
```


Your turn: Homogenized vs. full

Question. Is the difference between the two curves visible by eye? By what factor is the homogenized model faster and does the agreement survive a harsher, longer insult?

Hints. Read the printed mass-ratio gap and the two runtimes. The homogenized model should track the full one closely wherever the loading is slow and gradual.

✓ Answer: Homogenized vs. full

The homogenized curve overlays the full CMM closely at a fraction of the runtime ($O(N)$ vs. $O(N^2)$). **Learning:** it reproduces a **similar time trajectory** to the full model, cheaply.



Equilibrated CMM

Directly Solve for the End State

Set every rate to zero and solve the **algebraic** equilibrium, one root-find instead of thousands of time steps.

3D the 3D equilibrium conditions

$$\underbrace{\Upsilon^\alpha = 1 \Rightarrow \bar{\sigma}^* = \bar{\sigma}_h}_{\text{(A) growth balance}}$$

$$\underbrace{\mathbf{F}_n^\alpha = \text{const} \Rightarrow \sigma^\alpha = \sigma_h^\alpha}_{\text{(B) remodeling completed}}$$

$$\underbrace{\nabla \cdot \boldsymbol{\sigma} = \mathbf{0}}_{\text{(C) mechanical equilib.}}$$

0D scalar conditions at the evolved geometry

$$\bar{\sigma}^* = \bar{\sigma}_h, \quad \lambda_e^\alpha = G^\alpha, \quad \sigma_\theta = \frac{Pa}{h}.$$

Collapse to One Scalar Equation

Let the turnover constituents grow in fixed proportion (exact when they share gain and turnover). The 3D conditions (A)–(C) collapse to a *single* equation for the evolved stretch λ^* , consistent with the same λ , G^α , ϕ^α used throughout:

0D scalar equilibrium equation for λ^*

$$\phi^e \xi + \xi \frac{\bar{\sigma}_h - \sigma^e(G^e \lambda^*)}{\sigma_{\text{turn}} - \sigma_h^e} - \gamma (\lambda^*)^2 = 0 \quad (21)$$

with the Laplace exponent $n = 2$, γ the sustained load factor, ξ the surviving elastin fraction, ϕ^e the elastin mass fraction, $\sigma_h^e = \sigma^e(G^e)$ the elastin homeostatic stress, and σ_{turn} the homeostatic stress carried by the turnover constituents (collagen and smooth muscle). Once λ^* is known, the evolved mass ratio is $M/M_0 = \gamma(\lambda^*)^2$.

This equation *is* the fixed point of the homogenized ODEs ((19)–(20) with $d/dt = 0$), so its root coincides with the long-time limit of the transient theory.

The Punchline: Existence Sets the Capacity Limit

Equation (21) may have **no physical root**. Then no adapted state exists and the transient theories dilate without bound. So the solver is the cleanest **existence test**: no root $\Rightarrow$ the wall lacks the *capacity* to adapt. Existence is *necessary* for stability; whether the transient actually *reaches* an equilibrium that exists is a separate, gain-dependent question (dynamic stability). Near the boundary the transient slows critically, which is exactly why the instant algebraic test is so useful.

⚙️ Setup: Example (the equilibrated stability test, ex04)

Example. Skip the transient: solve the single algebraic equation (21) for the *final* adapted stretch λ^* . A physical root means the tissue adapts; *no root* means it dilates without bound. Scanning elastin loss then traces the stability boundary.

Run it.

```
uv run python run.py configs/ex04_equilibrated.yaml
```

Here `study: equilibrium` solves once for the given insult (and prints a transient cross-check: the homogenized plateau lands on λ^*), then the `scan: block` sweeps `elastin_surviving` and `report_critical: elastin` prints the critical surviving fraction ξ_{crit} .

Parametric study. Move the boundary by changing the wall's load-bearing *capacity*, then re-read ξ_{crit} :

```
... -set model.constituents.collagen.law.c1=600
... -set model.constituents.collagen.G=1.15
```



(Raising the gain leaves ξ_{crit} unchanged: gain sets how *fast* the transient adapts, not *whether* an equilibrium exists.)

Your turn: Equilibrated

Question. Does the algebraic λ^* land on the transient plateau where a root exists? What is the critical elastin loss beyond which *no* equilibrium exists at all?


✓ Answer: Equilibrated

Where a root exists, the equilibrated solve lands exactly on the transient end-state. Below ξ_{crit} there is *no* root: the transient dilates without bound.



Stability

Predict the Path to Heart Failure as a Stability Problem?



⚙️ Setup: Example (the stable/unstable map, ex05)

Example. A tissue *adapts* to an insult iff a mechanobiological equilibrium exists (21). Sweeping that instant test over (surviving elastin) $\times$ (pressure) draws the stable/unstable map. Then change *one* tissue property and watch the unstable “aneurysm” region shrink.

Run it.

```
uv run python run.py configs/ex05_stability.yaml
```

Here `study: map` builds the 2D grid over the two axes: via the `adapts` existence test and shades `adapt` (blue) vs. `runaway` (red); it also prints the fraction of the grid that adapts.

Parametric study. Enlarge the *adapts* region by giving the wall more load-bearing capacity, then compare the printed fraction:

```
... -set model.constituents.collagen.law.c1=600
... -set model.constituents.collagen.G=1.15
```

(Raising the gain does *not* move this map: existence is gain-independent. Gain governs *dynamic* stability instead, in the transient.)



Your turn: Stability map

Question. Which capacity change (stiffer collagen or larger deposition stretch) enlarges the stable region the most, and why does raising the *gain* leave the map unchanged? Elastin cannot be regrown in adults: what does that imply for aneurysm?

Hints. Compare the printed “fraction that adapts”. Stiffness and deposition stretch move the existence boundary; the gain does not, it sets how fast (and whether) the transient reaches an equilibrium that exists.


✓ Answer: Stability map

Stiffer collagen or a larger deposition stretch enlarge the *adapts* region: the wall carries more load, so it tolerates more elastin loss. The *gain* leaves this existence map unchanged (at equilibrium, production balances removal and the gain cancels). It instead sets *dynamic* stability, whether the transient actually reaches an equilibrium that exists (too weak adaptation runs away even when a root exists).




Future Perspectives


Clinical Application: Reversible Graft Stenosis



Systems-Biology Coupling: Signaling to Remodeling



Fluid–Solid–Growth: Closing the Hemodynamic Loop




3


FSGe: Long-Term Evolved Aneurysmal Equilibrium

Solid with $\tau_w \sim 1/r^3$


Wall shear
stress




**Fluid-Solid-Growth
interaction**



Velocity



Pressure



Open-Source Implementations

Under the Hood, What Each File Computes (1/2: the core)

The `gr` package is layered: a shared **continuum-mechanics core** (read, don't edit), then one file per G&R theory.

File	Computes	Continuum mechanics / G&R
<code>mechanics.py</code>	strain energy $W(\lambda_e)$, 2nd PK stress $S = dW/dE_e$, Cauchy $\sigma = \lambda_e^2 S$ (neo-Hookean, Fung)	finite-strain hyperelasticity of each constituent, in reference measures
<code>parameters.py</code>	Model/Constituent $(\phi_0, G^\alpha, k_d, K_\sigma)$; set-point $\sigma_h^\alpha = \sigma^\alpha(G^\alpha)$; Insult	deposition prestretch & homeo-static set-points ; the two insults
<code>geometry.py</code>	required Laplace stress $\bar{\sigma}_h \propto \lambda^2/m$; solves $\bar{\sigma}(\lambda) = \sigma_{\text{req}}$ for λ	mechanical equilibrium ; no root $\Rightarrow$ runaway
<code>history.py</code>	Result: $t, \lambda, M/M_0, \bar{\sigma}, \dots$	shared time-history output

Under the Hood, What Each File Computes (2/2: the theories)

File	Computes	G&R theory / role
kinematic_growth.py	split $\mathbf{F} = \mathbf{F}_e \mathbf{F}_g$, scalar θ ; $\dot{\theta} = k_g \theta (\bar{\sigma} / \bar{\sigma}_h - 1)$	Theory 1: kinematic growth
constrained_mixture.py	heredity integrals over cohorts; mass balance + mixture stress ($O(N^2)$)	Theory 2: full CMM
homogenized_cmm.py	two ODEs per constituent: $\dot{M}^\alpha, \dot{\lambda}_n^\alpha$ ($O(N)$)	Theory 3: homogenized CMM
equilibrated_cmm.py	rates $\rightarrow 0$: one algebraic solve for λ^*	Theory 4: equilibrated CMM
stability.py	adapts (does a root exist?), critical_elastin_loss	stability = existence of an equilibrium
scenario.py, configs/ run.py,	build model/geometry/insult from YAML; run a study; sweep	the exercise driver (input files)

Common thread: each theory changes only *how the supplied stress $\bar{\sigma}(\lambda)$ is built*, the same constituent laws and equilibrium solve are reused, so every difference you saw is the G&R law alone.

Research-Grade Open-Source Solvers

Project	What it does	Stack / group
svGrowth	constrained-mixture G&R of vessels (biaxial thin-wall); Stanford CBCL	
svFSGe	fast, strongly-coupled 3D fluid–solid–growth driver (the FSGe method)	Python + svFSI/svMultiPhysics; Yale (Pfaller lab)
FEMBE_Plugin	mechanobiologically equilibrated constrained-mixture material as a FEBio plugin	C++ / FEBio; Yale (Humphrey lab)
FEFSG_Plugin	fluid–solid–growth coupling as a FEBio plugin (thoracic-aorta hypertension & aneurysm cases)	C++ / FEBio; Yale (Humphrey lab)

`github.com/StanfordCBCL/svGrowth` `github.com/StanfordCBCL/svFSGe`

`github.com/yale-humphrey-lab/FEMBE_Plugin` `github.com/yale-humphrey-lab/FEFSG_Plugin`

How the Four Theories Answer: The Takeaways

Theory	Stability?	What it says
Kinematic growth	<i>Imposed</i>	Either reaches the <i>a-priori prescribed</i> homeostasis, or is unstable, nothing in between. Built in, not learned.
Constrained mixture (full)	<i>Predicted</i>	Adapts or dilates depending on the insult , the faithful reference.
Homogenized CMM	<i>Predicted</i>	A similar time trajectory to the full model, at a fraction of the cost.
Equilibrated CMM	<i>Instantly</i>	Matches the transient theories <i>if</i> an equilibrium exists; <i>no</i> root $\Rightarrow$ it does not match, unbounded growth.

Summary: What to Take Away

The biology. Living tissue continuously **turns over**; cells deposit new material *pre-stretched* at G^α , chasing a homeostatic mechanical set-point σ_h .

The four theories, one set-point, one question: does it adapt or run away?

- **Kinematic growth:** prescribe F_g ; *imposes* the set-point. Outcome is binary: reach the prescribed homeostasis, or go unstable.
- **Constrained mixture:** track every cohort; stability is *predicted* and **depends on the insult** (mild $\rightarrow$ adapt, severe $\rightarrow$ runaway).
- **Homogenized:** one mean natural configuration per constituent; a **similar time trajectory** to the full model at $O(N)$ not $O(N^2)$.
- **Equilibrated:** solve for the end state directly; **matches** the transient theories *iff* an equilibrium exists, and flags instability by finding no root.

Where it goes. These tools now drive **clinical prediction** (TEVGs), **systems-biology coupling**, and **fluid–solid–growth**.

Key references I

- Rodriguez, Hoger, McCulloch. Stress-dependent finite growth in soft elastic tissues. *J Biomech* (1994).
- Humphrey, Rajagopal. A constrained mixture model for growth and remodeling of soft tissues. *Math Models Methods Appl Sci* (2002).
- Cyron, Wilson, Humphrey. Mechanobiological stability. *J R Soc Interface* (2014).
- Cyron, Aydin, Humphrey. A homogenized constrained mixture (and mechanical analog) model. *Biomech Model Mechanobiol* (2016).
- Latorre, Humphrey. A mechanobiologically equilibrated constrained mixture model. *ZAMM* (2018).
- Latorre, Humphrey. Fast, rate-independent finite-element implementation of a 3D constrained mixture model. *Comput Methods Appl Mech Eng* (2020).
- Humphrey. Constrained mixture models of soft tissue growth and remodeling. *J Elasticity* (2021).
- Irons, Humphrey. Cell signaling model for arterial mechanobiology. *PLoS Comput Biol* (2020).
- Irons, Latorre, Humphrey. From transcript to tissue: multiscale modeling. *Ann Biomed Eng* (2021).
- Drews, Pepper, Best, . . . , Humphrey, Breuer. Spontaneous reversal of stenosis in tissue-engineered vascular grafts. *Sci Transl Med* (2020).

Key references II

- Figueroa, Baek, Taylor, Humphrey. A computational framework for fluid–solid–growth modeling. *Comput Methods Appl Mech Eng* (2009).
- Schwarz, Pfaller, . . . , Marsden. A fluid–solid–growth solver for cardiovascular modeling. *Comput Methods Appl Mech Eng* (2023).
- Pfaller, Latorre, . . . , Marsden. FSGe: a fast, strongly coupled 3D fluid–solid–growth method. *Comput Methods Appl Mech Eng* (2024).
- Humphrey, Eberth, Dye, Gleason. Fundamental role of axial stress in arterial mechanics. *J Biomech* (2009).
- Ambrosi, Ben Amar, Cyron, . . . , Humphrey, Kuhl. Growth and remodelling of living tissues. *J R Soc Interface* (2019).

Key references III

Code

- Seminar code (this deck's figures): `github.com/YaleCBL-Teaching/Growth-Remodeling`
- `github.com/StanfordCBCL/svGrowth`, `svFSGe`
- `github.com/yale-humphrey-lab/FEMBE_Plugin`, `FEFSG_Plugin`

Figure sources. Result figures reproduced with the `gr` package; literature figures cited per-slide.